

# Class10: Halloween mini project

Laura Sun (PID: A17923552)

## Table of contents

Quick overview of the dataset . . . . .	3
Winpercent and Pricepercent . . . . .	14
Exploring the correlation structure . . . . .	16
Principal Component Analysis . . . . .	16

As it is nearly Halloween and the half way point in the quarter, let's do a mini project to help us figure out the best candy!

Our data come from the 538 website and is available as a CSV file:

```
candy <- read.csv("candy-data.csv", row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crisped	rice	wafers
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent			
100 Grand	0	1	0	0.732	0.860	66.97173			
3 Musketeers	0	1	0	0.604	0.511	67.60294			
One dime	0	0	0	0.011	0.116	32.26109			
One quarter	0	0	0	0.011	0.511	46.11650			
Air Heads	0	0	0	0.906	0.511	52.34146			
Almond Joy	0	1	0	0.465	0.767	50.34755			

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

or

```
candy |>  
  nrow()
```

```
[1] 85
```

or call the tidyverse package!

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --  
v dplyr     1.1.4     v readr      2.1.5  
vforcats    1.0.1     v stringr    1.5.2  
v ggplot2   4.0.0     v tibble     3.3.0  
v lubridate 1.9.4     v tidyr     1.3.1  
v purrr     1.1.0  
-- Conflicts ----- tidyverse_conflicts() --  
x dplyr::filter() masks stats::filter()  
x dplyr::lag()   masks stats::lag()  
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
candy%>%  
  nrow()
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

My favorite candy is Hershey's Kisses, and the winpercent value is 55.37545.

```
candy["Hershey's Kisses", ]$winpercent
```

```
[1] 55.37545
```

Q4. What is the winpercent value for “Kit Kat”?

With tidyverse!

```
library(dplyr)
candy |> filter(rownames(candy)=="Kit Kat") |> select(winpercent)
```

```
winpercent
Kit Kat    76.7686
```

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

## Quick overview of the dataset

```
library(skimr)
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

**Variable type: numeric**

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

```
## skimr::skim(candy) to be easier
```

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

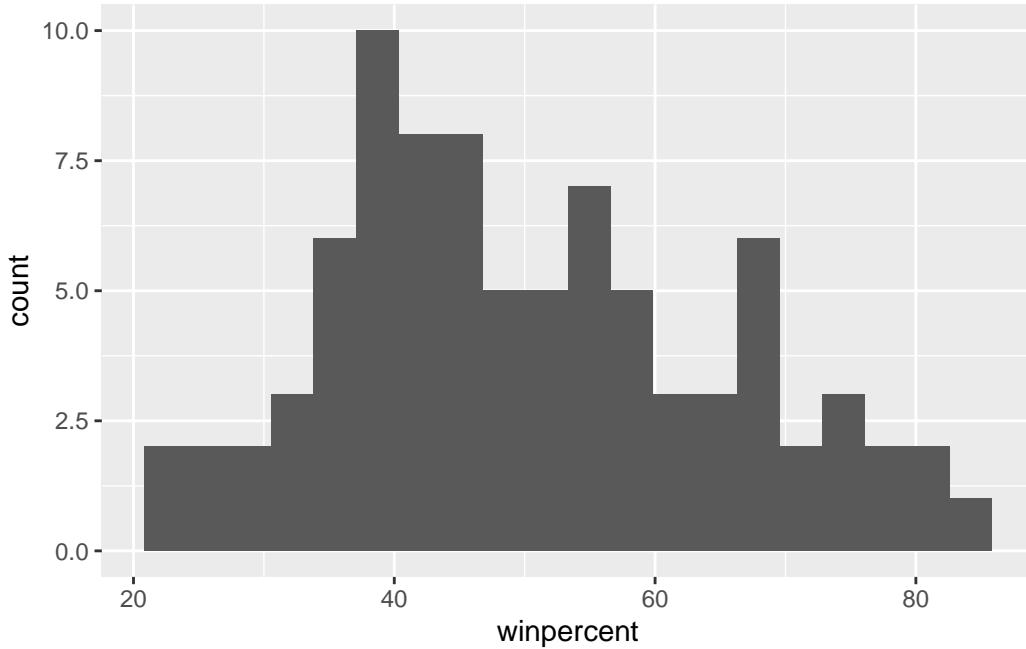
The winpercent is on a 0-100 scale. The rest are on 0-1 scale.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

The candy does not contain chocolate.

Q8. Plot a histogram of winpercent values

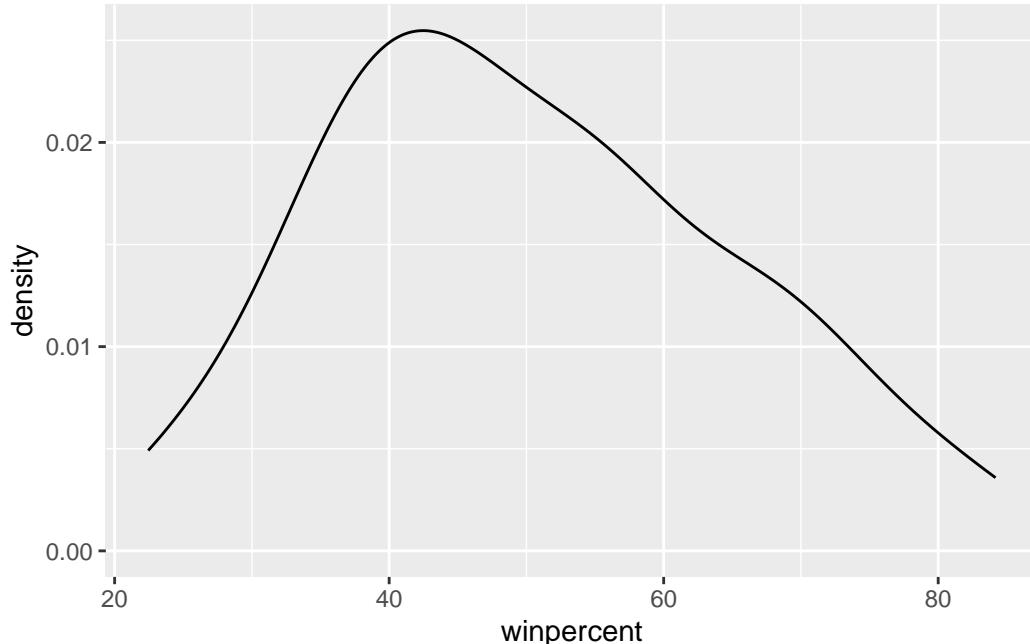
```
library(ggplot2)
ggplot(candy, aes(winpercent)) + geom_histogram(bins=20)
```



Q9. Is the distribution of winpercent values symmetrical?

It's not symmetrical. The graph is right skewed.

```
ggplot(candy, aes(winpercent)) + geom_density()
```



Q10. Is the center of the distribution above or below 50%?

```
mean(candy$winpercent)
```

```
[1] 50.31676
```

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

This depends on checking mean or median for center. Mean is above 50%, median is lower than 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
# 1. Find all chocolate candy in the dataset.  
# 2. Find their winpercent values  
# 3. Calculate the mean of these values.  
  
#4-6 Do the same for fruity candy.
```

```
# 7. Compare the mean winpercents of chocolate vs. fruity
# 8. Pick the higher one as winner.
```

```
choc.ind <- candy$chocolate==1
choc.win <- candy[choc.ind,]$winpercent
choc.mean <- mean(choc.win)
choc.mean
```

```
[1] 60.92153
```

```
fr.ind <- candy$fruity==1
fr.win <- candy[fr.ind,]$winpercent
fr.mean <- mean(fr.win)
fr.mean
```

```
[1] 44.11974
```

This is harder to read:

```
mean(candy[candy$chocolate==1,]$winpercent)
```

```
[1] 60.92153
```

```
mean(candy[candy$fruity==1,]$winpercent)
```

```
[1] 44.11974
```

Another way:

```
candy |> filter(chocolate==1) |> select(winpercent)
```

	winpercent
100 Grand	66.97173
3 Musketeers	67.60294
Almond Joy	50.34755
Baby Ruth	56.91455
Charleston Chew	38.97504
Hershey's Kisses	55.37545

Hershey's Krackel	62.28448
Hershey's Milk Chocolate	56.49050
Hershey's Special Dark	59.23612
Junior Mints	57.21925
Kit Kat	76.76860
Peanut butter M&M's	71.46505
M&M's	66.57458
Milk Duds	55.06407
Milky Way	73.09956
Milky Way Midnight	60.80070
Milky Way Simply Caramel	64.35334
Mounds	47.82975
Mr Good Bar	54.52645
Nestle Butterfinger	70.73564
Nestle Crunch	66.47068
Peanut M&Ms	69.48379
Reese's Miniatures	81.86626
Reese's Peanut Butter cup	84.18029
Reese's pieces	73.43499
Reese's stuffed with pieces	72.88790
Rolo	65.71629
Sixlets	34.72200
Nestle Smarties	37.88719
Snickers	76.67378
Snickers Crisper	59.52925
Tootsie Pop	48.98265
Tootsie Roll Juniors	43.06890
Tootsie Roll Midgies	45.73675
Tootsie Roll Snack Bars	49.65350
Twix	81.64291
Whoppers	49.52411

```
choc.mean>fr.mean
```

[1] TRUE

On average, chocolate candy (60.92153) is higher ranked than fruit candy (44.11974).

Q12. Is this difference statistically significant?

```
t.test(choc.win, fr.win)
```

### Welch Two Sample t-test

```
data: choc.win and fr.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

The difference is statistically significant. We have a p-value of 2.871e-08, which is a lot smaller than 0.05. We have enough evidence to support that there is a significant difference between chocolate and fruity candy's winpercent.

Q13. What are the five least liked candy types in this set?

```
candy |> arrange(winpercent) |> head(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0
	crisped	rice	wafer	hard	bar	pluribus
Nik L Nip	0	0	0		1	0.197
Boston Baked Beans	0	0	0		1	0.313
Chiclets	0	0	0		1	0.046
Super Bubble	0	0	0		0	0.162
Jawbusters	0	1	0		1	0.093
	percent	win				
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

or baseR:

```
position <- order(candy$winpercent)
head(candy[position,], 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0
	crispedrice	wafer	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

Q14. What are the top 5 all time favorite candy types out of this set?

```
candy |> arrange(desc(winpercent)) |> head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
	crispedrice	wafer	hard	bar	pluribus	sugarpercent
Reese's Peanut Butter cup	0	0	0	0	0	0.720
Reese's Miniatures	0	0	0	0	0	0.034
Twix	1	0	1	0	0	0.546
Kit Kat	1	0	1	0	0	0.313
Snickers	0	0	1	0	0	0.546
	pricepercent	winpercent				

Reese's Peanut Butter cup	0.651	84.18029
Reese's Miniatures	0.279	81.86626
Twix	0.906	81.64291
Kit Kat	0.511	76.76860
Snickers	0.651	76.67378

or

```
candy |> arrange(winpercent) |> tail(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crispedrice	wafer	hard	bar	pluribus	sugarpercent
Snickers	0	0	1		0	0.546
Kit Kat	1	0	1		0	0.313
Twix	1	0	1		0	0.546
Reese's Miniatures	0	0	0		0	0.034
Reese's Peanut Butter cup	0	0	0		0	0.720
	pricepercent	winpercent				
Snickers	0.651	76.67378				
Kit Kat	0.511	76.76860				
Twix	0.906	81.64291				
Reese's Miniatures	0.279	81.86626				
Reese's Peanut Butter cup	0.651	84.18029				

or

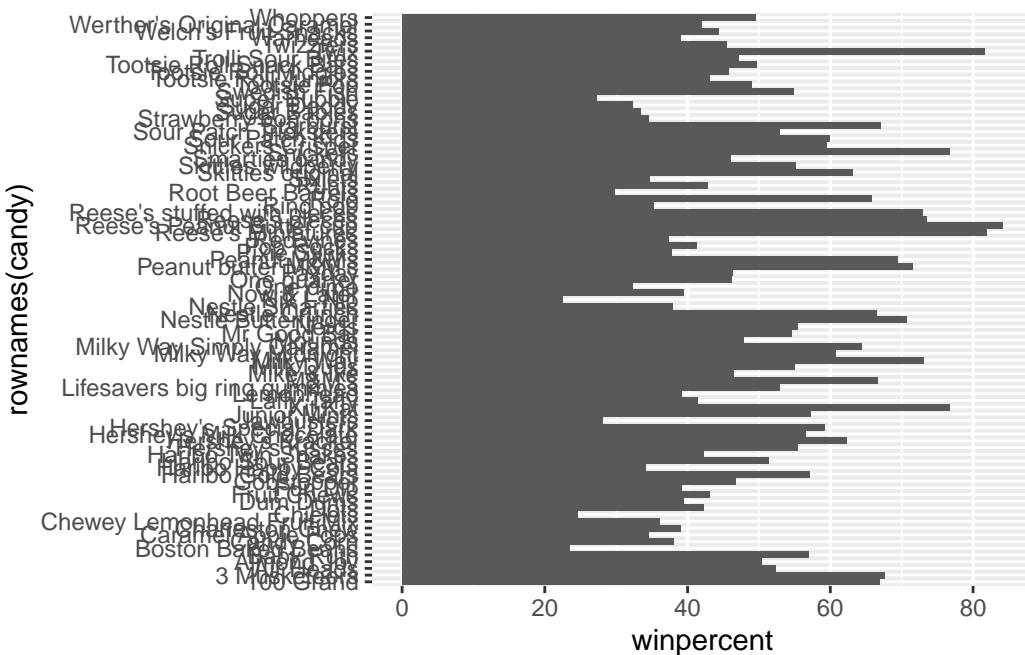
```
candy |> arrange(-winpercent) |> head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1
	crispedrice	wafer	hard	bar	pluribus	sugarpercent
Reese's Peanut Butter cup	0	0	0		0	0.720

	pricel	percent	winpercent		
Reese's Miniatures	0	0	0	0.034	
Twix	1	0	1	0	0.546
Kit Kat	1	0	1	0	0.313
Snickers	0	0	1	0	0.546
Reese's Peanut Butter cup	0.651	84.18029			
Reese's Miniatures	0.279	81.86626			
Twix	0.906	81.64291			
Kit Kat	0.511	76.76860			
Snickers	0.651	76.67378			

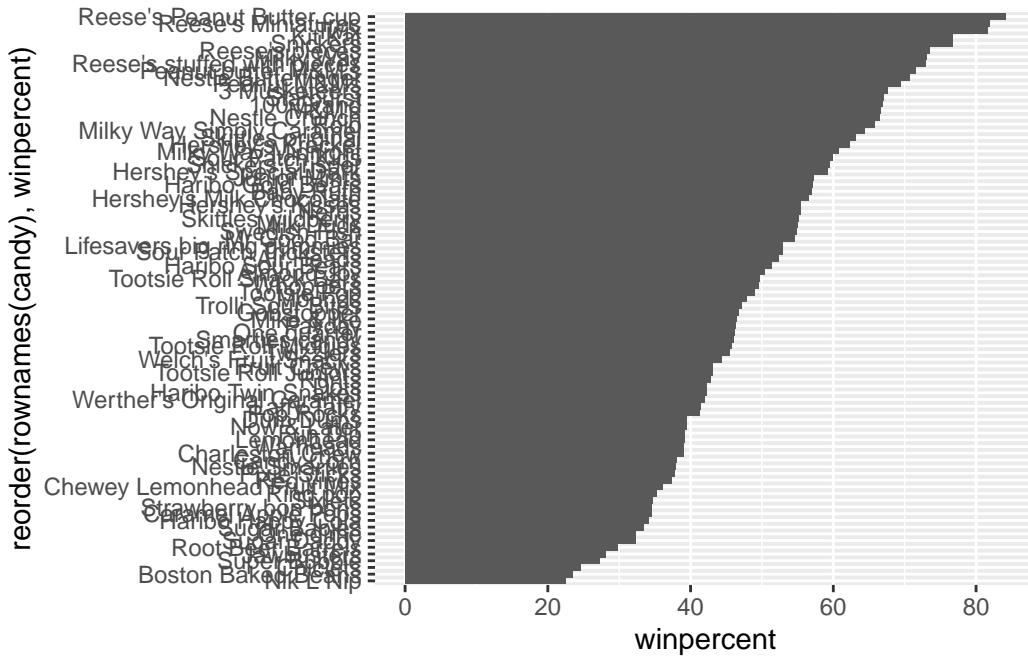
Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy, aes(winpercent, rownames(candy))) + geom_col()
```



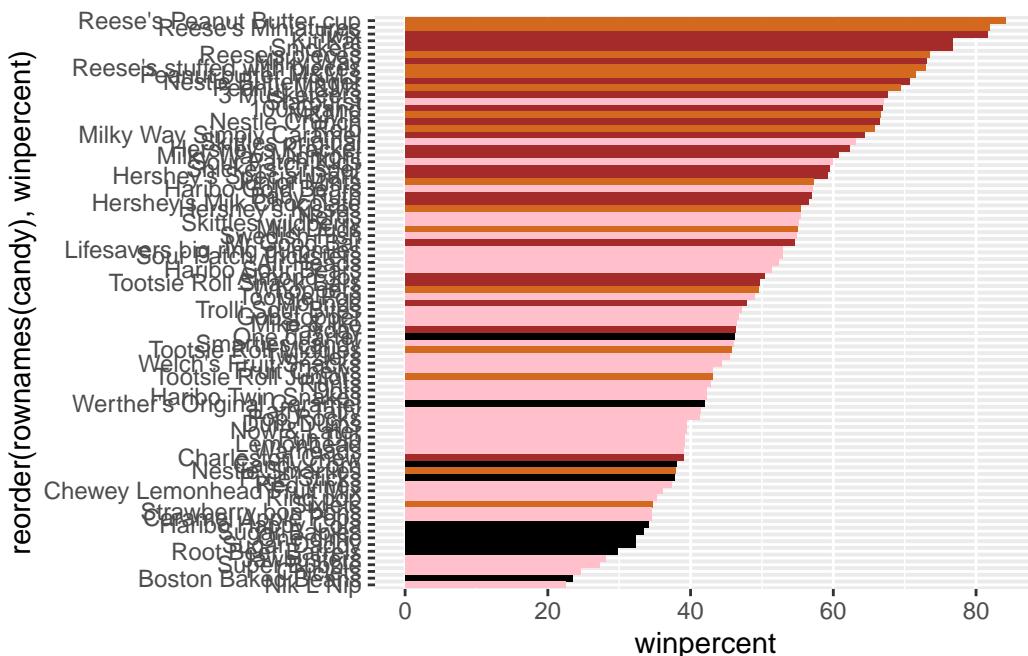
Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) + aes(x=winpercent, y=reorder(rownames(candy), winpercent)) + geom_col()
```



Adding colors based on the type of candy!

```
my_cols <- rep("black", nrow(candy)) # returns a vector of "black", if used in graph, all black
my_cols[as.logical(candy$chocolate)] <- "chocolate"
my_cols[as.logical(candy$fruity)] <- "pink"
my_cols[as.logical(candy$bar)] <- "brown" #for chocolate bar, overide chocolate color with brown
ggplot(candy) + aes(x=winpercent, y=reorder(rownames(candy), winpercent)) + geom_col(fill=my_cols)
```



Q17. What is the worst ranked chocolate candy?

Sixlets is the worst ranked chocolate candy.

Q18. What is the best ranked fruity candy?

Starburst is the best ranked fruity candy.

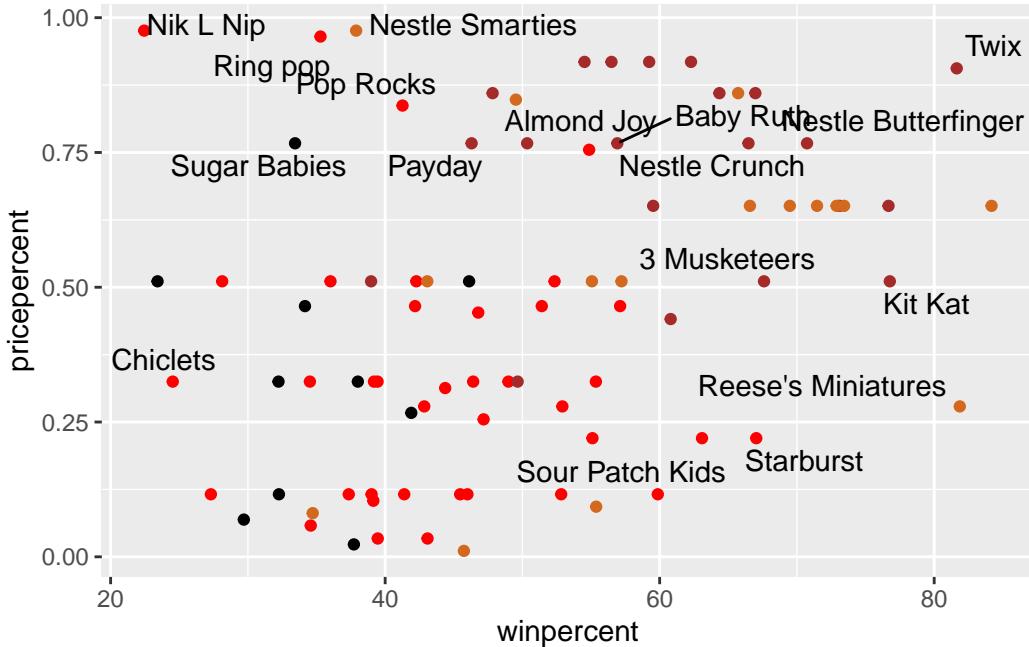
## Winpercent and Pricepercent

A plot with both variables/columns winpercent and pricepercent.

```
library(ggrepel)
my_cols[as.logical(candy$fruity)] <- "red"

ggplot(candy) + aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(max.overlaps=7)
```

Warning: ggrepel: 68 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Based on the graph, Reese's Miniatures seem to be highly popular and affordable.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent)
tail(candy[ord,c(11,12)], n=5)
```

	pricepercent	winpercent
Hershey's Special Dark	0.918	59.23612
Mr Good Bar	0.918	54.52645
Ring pop	0.965	35.29076
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719

The five most expensive candies are listed above. Among them, Nik L Nip is the least popular (with a winpercent of around 22)

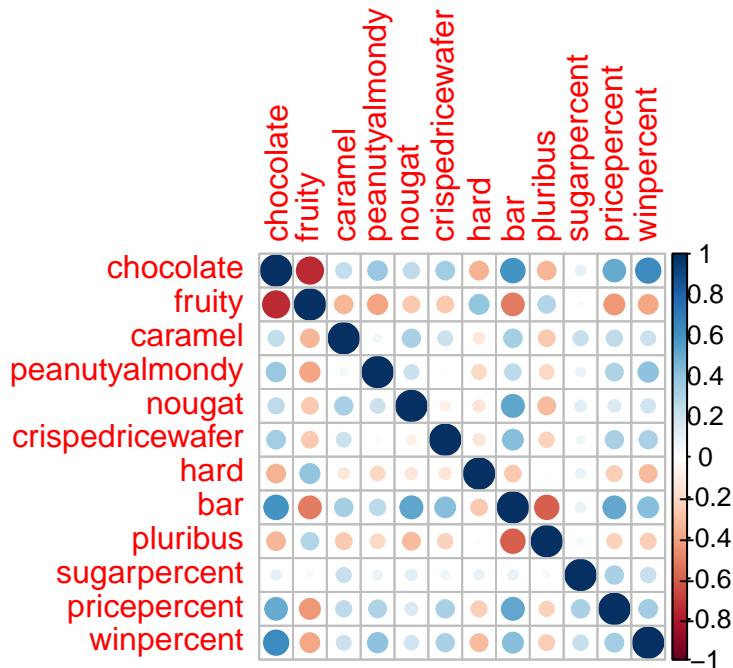
## Exploring the correlation structure

Now that we've explored the dataset a little, we'll see how the variables interact with one another. We'll use correlation and view the results with the corrplot package to plot a correlation matrix.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



## Principal Component Analysis

The function to use is called `prcomp()` with an optional `scale=T/F` argument.

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```

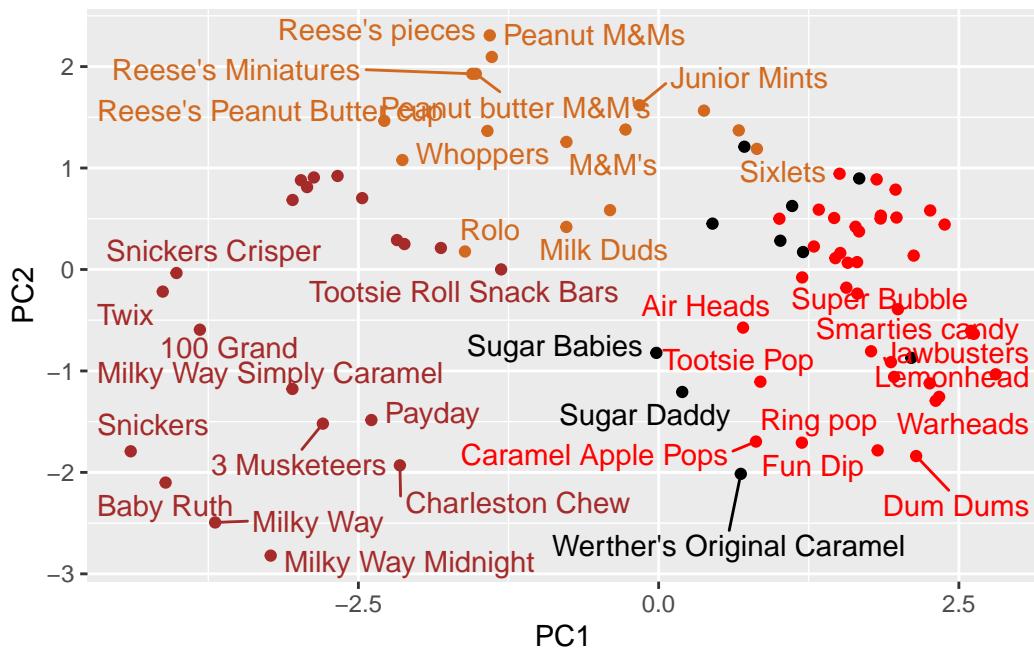
Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

Our main pca result figure:

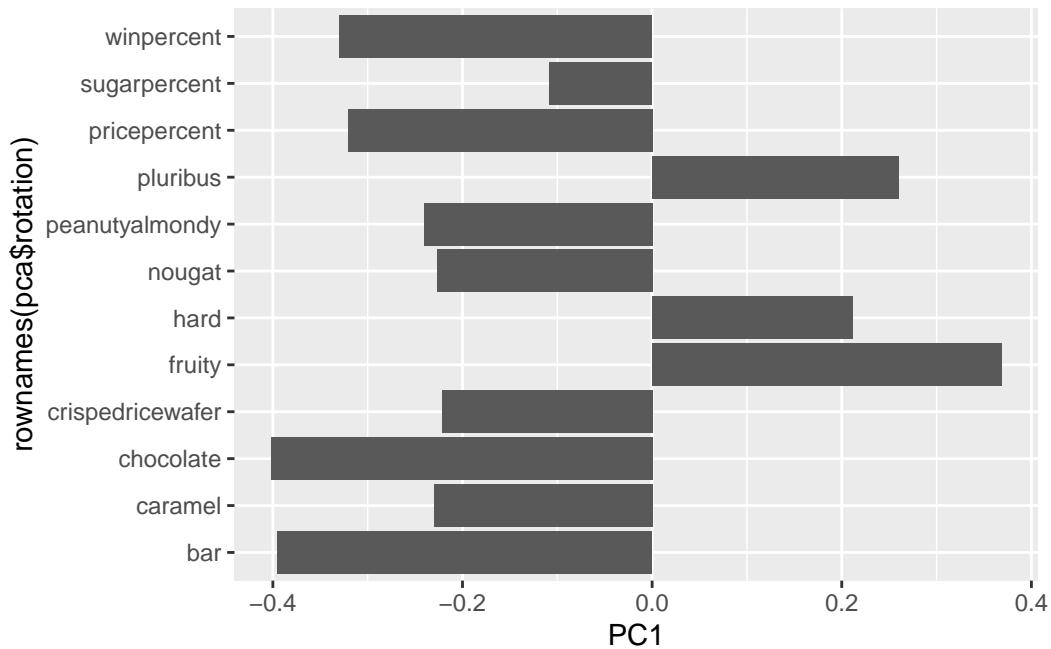
```
ggplot(pca$x) + aes(PC1, PC2, label=rownames(pca$x)) + geom_point(col=my_cols) + geom_text_r
```

Warning: ggrepel: 48 unlabeled data points (too many overlaps). Consider increasing max.overlaps



We should also examine the “variable loadings” or contributions of the original variables to the new PCs.

```
ggplot(pca$rotation) + aes(PC1, rownames(pca$rotation)) + geom_col()
```



```
p <- ggplot(pca$x) + aes(PC1, PC2, label=rownames(pca$x)) + geom_point(col=my_cols) + geom_t
```

Interactive plots than can be zoomed on and “brushed” over can be made with the **plotly** package. It’s output is interactive and will not render to PDF :(